

EE 156: Advanced Topics in Computer Architecture

Spring 2023
Tufts University

Instructor: Prof. Mark Hempstead
mark.hempstead@tufts.edu

Lecture 1: Bandwidth, latency, power and
Performance Metrics
[Chapter 1]

EE 156/CS140 Mark Hempstead

1

Lecture Outline

- Technology Trends
- Bandwidth vs. Latency
- Performance Metrics
 - Averaging
 - Amdahl's Law
 - CPU performance equation
- Benchmarks and Quantitative Analysis

Homework: **Read Chapter 1** for more detail

EE194/CS140 Mark Hempstead


2

Technology scaling

- Everyone has heard of Moore's Law. It's probably been mentioned in most newspapers at some point. But what does it really mean?
- Basic premise: every generation (1-2 years) all dimensions scale by 0.7.
 - $.7^2=0.5$, so devices/unit area doubles
 - This is less of a natural law, and more of a prediction of human progress. Actually making this happen has huge consequences

EE194/CS140 Mark Hempstead

3



Technology Scaling Road Map (ITRS no IRDS)

Year	2004	2006	2008	2010	2012	2014	2017	2019	2021	2024
Feature size (nm)	90	65	45	32	22	14	"10"	"7"	5	"3"
Intg. Capacity (BT)	2	4	8	16	33	83	162	332	651	1808

- Fun facts about 45nm transistors
 - 30 million can fit on the head of a pin
 - You could fit more than 2,000 across the width of a human hair
 - If car prices had fallen at the same rate as the price of a single transistor has since 1968, a new car today would cost about 1 cent

5

- ### We've done clever things
- Architectural Innovations
 - Massive pipelining (good and bad!)
 - Branch Prediction, Register Renaming, OOO-issue/execution, Speculation
 - Hierarchical caches and prefetching
 - Virtual memory acceleration (TLBs)
 - Circuit/Logic Innovations
 - New logic circuit families (dynamic logic)
 - Better CAD tools
 - Advanced computer arithmetic
- EE194/CS140 Mark Hempstead
- 6

How have we used these transistors?


- More functionality on one chip
 - Early 1980s – 32-bit microprocessors
 - Late 1980s – On Chip Level 1 Caches
 - Early/Mid 1990s – 64-bit microprocessors, superscalar (ILP)
 - Late 1990s – On Chip Level 2 Caches
 - Early 2000s – Chip Multiprocessors, On Chip Level 3 Caches
 - Early 2010s – Many-Core, SoC integration, specialized hardware
 - Late 2010s – Many-Core Mobile, Domain Specific Processors (e.g. Google TPU), Support for Virtualization
- What is next?
 - How much more cache can we put on a chip? (Itanium2)
 - How many more cores can we put on a chip? (Niagara, etc)
 - What else can we put on chips? (Apple and Mobile SoCs)
 - What can we build custom chips for? (Machine Learning (TPU))

EE 126 Mark Hempstead

7


Another Example of Moore's Law Impact


DRAM capacity growth over 3 decades



8

But What Happened to Clock Rates and Why?






Bandwidth and Latency

- What is the difference between bandwidth & latency?
- Not every performance measure has benefited equally from scaling
- Bandwidth (or throughput)
 - Total work done in a given time
 - 10,000-25,000X improvement for processors
 - 300-1200X improvement for memory and disks
- Latency (or response time)
 - Time between start and completion of an event
 - 30-80X improvement for processors
 - 6-8X improvement for memory and disks

EE194/CS140 Mark Hempstead



Rule of Thumb for Latency Lagging BW

- In the time that bandwidth doubles, latency improves by no more than a factor of 1.2 to 1.4
- Stated alternatively:
Bandwidth improves by more than the square of the improvement in Latency

EE194/CS140 Mark Hempstead

13

Why has latency lagged?

- When we halve delay (and double frequency), we halve latency and double bandwidth. Both benefit. The same is true for most things that benefit latency (e.g., faster disks).
- Adding more transistors rarely benefits both equally:
 - Widen datapaths, add more cores or execution units, etc. all directly benefit bandwidth but not latency
 - Using the transistors to build a faster multiplier (e.g., a Wallace tree) or adder (e.g., better carry lookahead) would help latency; but we are far into diminishing returns in these areas.
- Most algorithms we will talk about (more pipelining, branch prediction, etc) help bandwidth and slightly *hurt* latency.
 - We'll keep discussing algorithms' effects on latency, bandwidth & power.

EE194/CS140 Mark Hempstead

14

Metrics

- We've seen that power, frequency, latency and bandwidth do not improve at the same rate.
- So what do we measure when we benchmark architectures against each other?
 - Most common answer: just measure execution time.
 - Execution time may measure latency (if we're only measuring the time for one action) or throughput (if we're measuring the time for 10K tasks and the program overlaps them).
 - Anyone can munge the results to, e.g., look only at computers in the same power class.

EE194/CS140 Mark Hempstead

15

Performance Metrics

- How do we decide the tasks?
- Some benchmarks
 - What the customer cares about, real applications
 - Representative programs (SPEC, SYSMARK, etc)
 - Kernels: Code fragments from real programs (Linpack)
 - Toy Programs: Quicksort
 - Synthetic Programs: Just a representative instruction mix (Whetsone, Dhrystone)

Better
↑

EE194/CS140 Mark Hempstead

16

Averaging over a suite of benchmarks

- We can average total execution time:

$$\frac{1}{n} \sum_{i=1}^n \text{Time}_i$$

- This is the *arithmetic* mean
 - It should be used when measuring performance in execution *times*

EE194/CS140 Mark Hempstead

17

Averaging over a suite of benchmarks

- We can average the “normalized execution time,” i.e., time / time-on-a-reference-machine
- Can only use **geometric mean** (arithmetic mean can vary depending on the reference machine)

$$\sqrt[n]{\prod_{i=1}^n \text{ExecutionTimeRatio}_i}$$

- This is what SPEC does
- Problem: Ratio not Execution Time is the result

EE194/CS140 Mark Hempstead

18

Amdahl's Law (Law of Diminishing Returns)

- Very Intuitive – Make the Common case fast

$$\text{Speedup} = \frac{\text{Execution Time for task without enhancement}}{\text{Execution Time for task using enhancement}}$$

$$\text{Execution time}_{\text{new}} = \text{Execution time}_{\text{old}} \times$$

$$\text{Overall Speedup} \rightarrow \left((1 - \text{Fraction}_{\text{enhanced}}) + \frac{\text{Fraction}_{\text{enhanced}}}{\text{Speedup}_{\text{enhanced}}} \right)$$

ECEC 621
Mark Hempstead

19

Amdahl's Law Example

Speed up 95% of the task by 1.1x:

$$\text{Speedup}_{\text{Overall}} = 1 / ((1 - .95) + (.95 / 1.1)) = 1.094$$

Speed up 5% of the task by 10x:

$$\text{Speedup}_{\text{Overall}} = 1 / ((1 - .05) + (.05 / 10)) = 1.047$$

Speed up 5% of the task infinitely:

$$\text{Speedup}_{\text{Overall}} = 1 / (1 - .05) = 1.052$$

Make the common
case fast!

EE194/CS140 Mark Hempstead

20

Measure MIPS?

- MIPS = instruction count / (execution time x 10⁶)
= clock rate / (CPI x 10⁶)
- Problems
 - ISAs are not equivalent, e.g. RISC vs. CISC
 - 1 CISC instruction may equal many RISC!
 - Programs use different instruction mixes
 - May be ok when comparing same benchmarks, same ISA, same compiler, same OS

EE194/CS140 Mark Hempstead

21

Measure MFLOPS

- Same as MIPS, just FP ops
- Not useful either
 - FP-intensive apps needed
 - Traditionally, FP ops were slow, INT can be ignored
 - BUT, now memory ops can be the slowest!
- “Peak MFLOPS” is a common marketing fallacy
 - Basically, it just says #FP-pipes X Clock Rate

Measure GHz?

- Is this a metric? Maybe as good as the others...
- One number, no benchmarks, what can be better?
- Many designs *are* frequency driven

Processor	Clock Rate	SPEC FP2000
IBM POWER3	450 MHz	434
Intel PIII	1.4 GHz	456
Intel Pentium 4	2.4 GHz	833
Itanium-2	1.0 GHz	1356

Performance: What to measure

- To increase predictability, collections of benchmark applications, called *benchmark suites*, are popular
- **SPEC_{CCPU}**: popular desktop benchmark suite
 - SPECint2000 has 12 integer, SPECfp2000 has 14 integer pgms
 - SPEC_{CCPU}2006 announced Spring 2006
 - SPEC_{FS} (NFS file server) and SPEC_{Web} (WebServer) added as server benchmarks
- **Transaction Processing Council** measures server performance and cost-performance for databases
 - **TPC-C** Complex query for Online Transaction Processing
 - TPC-H models ad hoc decision support
 - TPC-W a transactional web benchmark
 - TPC-App application server and web services benchmark
- **Splash2** Parallel Benchmark Suite (use in lab)
- New Emerging suites: PARSEC, RODINA, BIO-BENCH
- Embedded industry has its own benchmarks: Dhrystone, EEMBC, SPECjvm

New Emerging Benchmarks


- Significant updates over the few years with new application domains
- SPEC2017
 - Designed to provide performance measurements that can be used to compare compute-intensive workloads on different computer systems
 - 43 benchmarks organized into four suites: SPECspeed 2017 Integer, SPECspeed 2017 FP, SPECrate 2017 Integer, and SPECrate 2017 FP
- MLPerf <https://mlperf.org/>
 - A broad machine learning benchmark suite for measuring performance of ML software frameworks, ML hardware accelerators, and ML cloud platforms.
 - Image classification, Object Detection, Translation, Recommendation, Reinforcement Learning



EE 156/CS140 Mark Hempstead

25

What's a Clock Cycle?



- Old days: 10 levels of gates (e.g. 10 FO4)
- Today: determined by numerous time-of-flight issues + gate delays
 - clock propagation, wire lengths, drivers

EE194/CS140 Mark Hempstead

26

THE Performance Equation

$$\text{CPU time} = \text{Instruction_count} \times \text{CPI} \times \text{clock_cycle}$$

or

$$\text{CPU time} = \frac{\text{Instruction_count} \times \text{CPI}}{\text{clock_rate}}$$

□ These equations separate the three key factors that affect performance

- Can measure the CPU execution time by running the program
- The clock rate is usually given
- Can measure overall instruction count by using profilers/ simulators without knowing all of the implementation details
- CPI varies by instruction type and ISA implementation for which we must know the implementation details

27

CPU Performance Equation

- Execution Time = $\frac{\text{seconds}}{\text{program}}$

$$\frac{\text{instructions}}{\text{program}} \times \frac{\text{cycles}}{\text{instruction}} \times \frac{\text{seconds}}{\text{cycle}}$$

Program	Compiler (Scheduling)	Technology
Architecture (ISA)	Organization (uArch)	Physical Design
Compiler	Microarchitects	Circuit Designers

EE194/CS140 Mark Hempstead

28

Alternately put...

- The equation again for total time:

$$\frac{\text{instructions}}{\text{program}} \times \frac{\text{cycles}}{\text{instruction}} \times \frac{\text{seconds}}{\text{cycle}}$$

Can be re-phrased as

total time = (total instructions)*CPI/freq.

EE194/CS140 Mark Hempstead

29


More CPI

- CPI can be useful even when a program is a mix of instruction types, each with their own CPI:
- Total CPU cycles = $\sum_i IC_i CPI_i$, where the summation is over each different instruction type.
- This can tell us where to focus our efforts:
 - But always remember Amdahl's Law
 - And remember that improving one instruction may hurt another, or even hurt your cycle time. Many published studies do in fact forget this.

EE194/CS140 Mark Hempstead

30


Why Didn't This Happen?



[Taylor, DAC and DaSi 2012]

31

Why Power stopped increasing

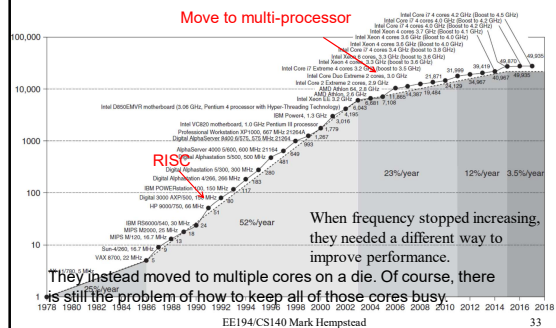


EE194/CS140 Mark Hempstead

32

Crossroads: Uniprocessor Performance

From Hennessy and Patterson, *Computer Architecture: A Quantitative Approach*, 6th edition, 2018



EE194/CS140 Mark Hempstead

33

Chapter 1 Summary

- Technology Trends have a significant impact on design
 - Scaling of wires, transistors, power
 - Latency has been lagging BW
 - Designers have plenty of transistors? Just don't know what to do with them.
- Performance should be measured quantitatively
 - However there are different benchmarks and metrics available

H&P Turing Lecture



Backup

- List of Benchmarks
- Additional material hidden in this PPT. It is discussed in chapter 1, but we will cover it later in the class.

SPEC CPU2006: Integer Benchmarks

400.perlbenc	C	Perl Programming Language
401.bzip2	C	Bzip compression
403.gcc	C	GCC Compiler
429.mcf	C	Optimization. Vehicle scheduling for public transit
445.gobmk	C	AI the game of GO
456.hmmcr	C	Protein sequence analysis using profile hidden
458.sjeng	C	AI Chess program
462.libquant	C	Simulates a quantum computer, running Shor's
464.h264ref	C	Video Compression, H.264/AVC standard
471.omnetpp	C++	Uses the OMNet++ discrete event simulator to
473.astar	C++	Pathfinding library for 2D maps, including the well
483.xalancb	C++	XML Processing

EE194/CS140 Mark Hempstead

37

SPEC CPU2006: Floating Point Benchmarks

410.bwaves	Fortran	Fluid Dynamics
416.gamess	Fortran	Quantum Chemistry
433.mlc	C	Physics / Quantum Chromodynamics
434.zeusmp	Fortran	Physics / computational fluid dynamics
435.gromacs	C, Fortran	equations of motion for hundreds to millions
436.cactusADM	C, Fortran	Fortran Physics / General Relativity
437.jeslie3d	Fortran	Large-Eddy Simulations
444.namd	C++	Biology / Molecular Dynamics
447.deall	C++	Finite Element Analysis
450.soplex	C++	Linear Programming, Optimization
453.povray	C++	Image Ray-tracing
454.calculx	C, Fortran	Structural Mechanics
459.GemsFDTD	Fortran	Computational Electromagnetics
465.tonto	Fortran	Quantum Chemistry
470.lbm	C	Fluid Dynamics
481.wrf	C, Fortran	Weather modeling
482.sphinx3	C	Speech recognition

EE194/CS140 Mark Hempstead

38

Parsec Benchmarks

- SPEC CPU does not include multithreaded and emerging workloads
- Other suites such as PARSEC, SPLASH-2, Rodina (CUDA) include such workloads
- List of PARSEC benchmarks:
 - blackscholes** - Option pricing with Black-Scholes Partial Differential Equation (PDE)
 - bodytrack** - Body tracking of a person
 - cannal** - Simulated cache-aware annealing to optimize routing cost of a chip design
 - dedup** - Next-generation compression with data deduplication
 - facesim** - Simulates the motions of a human face
 - ferret** - Content similarity search server
 - fluidanimate** - Fluid dynamics for animation purposes with Smoothed Particle Hydrodynamics (SPH) method
 - freqmine** - Frequent itemset mining
 - raytrace** - Real-time raytracing
 - streamcluster** - Online clustering of an input stream
 - swaptions** - Pricing of a portfolio of swaptions
 - vips** - Image processing (Project Website)
 - x264** - H.264 video encoding

EE194/CS140 Mark Hempstead

39

Server Benchmarks

- TPC-C (Online-Transaction Processing, OLTP)
 - Models a simple order-entry application
 - Thousands of concurrent database accesses
 - TPM = transactions per minute

System	# / Processor	tpm	\$/tpm
Fujitsu PrimePower	128, 563MHz SPARC64	455K	\$28.58
HP SuperDome	64, 875MHz PA8700	423K	\$15.64
IBM p690	32, 1300MHz POWER4	403K	\$17.80

- TPC-H (Ad-hoc, decision support)
 - Data warehouse, backend analysis tools for data

EE194/CS140 Mark Hempstead

40