

EE 156: Advanced Topics in Computer Architecture

Spring 2023
Tufts University

Instructor: Prof. Mark Hempstead
mark@ece.tufts.edu

Lecture 7: Branch Prediction and
Advanced Branch Prediction

EE156/CS140 Mark Hempstead

1

Outline

Unit 2: Microarchitecture and the Pipeline (3 weeks)

- *Instructions Introducing the RISC-V ISAs*
- *Basic Pipelining Review [Appendices A, C and K]*
- Hardware instruction-level parallelism (ILP) and Tomosulo's algorithm [Ch 3]
- **Advanced Branch Prediction**

EE156/CS140 Mark Hempstead

2

Branch prediction

- We looked at two strategies for speeding up branches; assume-taken and assume-not-taken. Neither worked well.
- Now we'll look at branch prediction.
 - Hardware guesses branch outcome
 - Start fetching from guessed address
 - Flush the pipe on a mis-predict.

EE156/CS140 Mark Hempstead

3

People branch-predict too

- Buy an engagement ring, predicting a “yes” answer
 - If you guessed wrong, then throw away the ring?
- Pick which way to drive to work based on what the traffic was like yesterday
 - If today’s backups are different, then make excuses when you get to work

EE156/CS140 Mark Hempstead

4

Branch prediction is hard

- Making predictions is hard, especially when they’re about things that haven’t happened yet.
 - Yogi Berra
- Corollary: branch prediction is hard

EE156/CS140 Mark Hempstead

5


Branch Prediction Performance

- Parameters
 - **Branch: 20%**, load: 20%, store: 10%, other: 50%
 - 75% of branches are taken
- Dynamic branch prediction
 - Branches predicted with 95% accuracy
 - $\text{CPI} = 1 + 20\% * 5\% * 2 = \mathbf{1.02}$
 - What is the CPI without branch prediction? (assume a cost of 2 cycles to resolve a taken branch and that 50 of branches are taken)
 - $\text{CPI} = 1 + 20\% * 50\% * 2 = \mathbf{1.2}$

EE156/CS140 Mark Hempstead

6

Dynamic Branch Prediction Components




- Step #1: is it a branch?
 - Easy after decode...
- Step #2: is the branch taken or not taken?
 - **Direction predictor** (applies to conditional branches only)
 - Predicts taken/not-taken
- Step #3: if the branch is taken, where does it go?
 - Branch **target predictor**
 - Easy after decode...

EE156/CS140 Mark Hempstead

7

Branch Prediction

- **Learn from past, predict the future**
 - Record the past in a hardware structure
- **Direction predictor (DIRP)**
 - Map conditional-branch PC to taken/not-taken (T/N) decision
 - Individual conditional branches often biased or weakly biased
 - 90%+ one way or the other considered **"biased"**
 - Why? Loop back edges, checking for uncommon conditions
- **Branch history table (BHT):** simplest predictor
 - PC indexes table of bits (0 = Not taken, 1 = Taken), no tags
 - Essentially: branch will go same way it went last time



EE156/CS140 Mark Hempstead

Branch History Table (BHT)

- **Branch history table (BHT):** simplest direction predictor
 - PC indexes table of bits (0 = N, 1 = T), no tags
 - Predicts the branch to go same way it went last time
 - Problem: **inner loop branch** below


```
for (i=0; i<100; i++)
  for (j=0; j<3; j++)
    // whatever
```

 - It will be wrong twice per inner loop; once each at the beginning and the end.
 - Branch predictor "changes its mind too quickly"

Time	State	Prediction	Outcome	Result?
1			T	
2			T	
3			T	
4			N	
5			T	
6			T	
7			T	
8			N	
9			T	
10			T	
11			T	
12			N	

EE156/CS140 Mark Hempstead

9

Branch History Table (BHT)

- **Branch history table (BHT):**
simplest direction predictor

- PC indexes table of bits (0 = N, 1 = T), no tags
- Predicts the branch to go same way it went last time
- Problem: **inner loop branch** below

```
for (i=0; i<100; i++)
  for (j=0; j<3; j++)
    // whatever
```
- It will be wrong twice per inner loop; once each at the beginning and the end.
- Branch predictor “changes its mind too quickly”

Time	State	Prediction	Outcome	Result?
1	N	N	T	Wrong
2	T	T	T	Correct
3				
4				
5				
6				
7				
8				
9				
10				
11				
12				

EE156/CS140 Mark Hempstead

10

Branch History Table (BHT)

- **Branch history table (BHT):**
simplest direction predictor

- PC indexes table of bits (0 = N, 1 = T), no tags
- Predicts the branch to go same way it went last time
- Problem: **inner loop branch** below

```
for (i=0; i<100; i++)
  for (j=0; j<3; j++)
    // whatever
```
- It will be wrong twice per inner loop; once each at the beginning and the end.
- Branch predictor “changes its mind too quickly”

Time	State	Prediction	Outcome	Result?
1	N	N	T	Wrong
2	T	T	T	Correct
3	T	T	T	Correct
4				
5				
6				
7				
8				
9				
10				
11				
12				

EE156/CS140 Mark Hempstead

11

Branch History Table (BHT)

- **Branch history table (BHT):**
simplest direction predictor

- PC indexes table of bits (0 = N, 1 = T), no tags
- Predicts the branch to go same way it went last time
- Problem: **inner loop branch** below

```
for (i=0; i<100; i++)
  for (j=0; j<3; j++)
    // whatever
```
- It will be wrong twice per inner loop; once each at the beginning and the end.
- Branch predictor “changes its mind too quickly”

Time	State	Prediction	Outcome	Result?
1	N	N	T	Wrong
2	T	T	T	Correct
3	T	T	T	Correct
4	T	T	T	Correct
5				
6				
7				
8				
9				
10				
11				
12				

EE156/CS140 Mark Hempstead

12

Branch History Table (BHT)

- **Branch history table (BHT):**
simplest direction predictor

- PC indexes table of bits (0 = N, 1 = T), no tags
- Predicts the branch to go same way it went last time
- Problem: **inner loop branch** below

```
for (i=0; i<100; i++)
  for (j=0; j<3; j++)
    // whatever
```
- It will be wrong twice per inner loop; once each at the beginning and the end.
- Branch predictor “changes its mind too quickly”

Time	State	Prediction	Outcome	Result?
1	N	N	T	Wrong
2	T	T	T	Correct
3	T	T	T	Correct
4	T	T	N	Wrong
5	N			
6				
7				
8				
9				
10				
11				
12				

EE156/CS140 Mark Hempstead

13

Branch History Table (BHT)

- **Branch history table (BHT):**
simplest direction predictor

- PC indexes table of bits (0 = N, 1 = T), no tags
- Predicts the branch to go same way it went last time
- Problem: **inner loop branch** below

```
for (i=0; i<100; i++)
  for (j=0; j<3; j++)
    // whatever
```
- It will be wrong twice per inner loop; once each at the beginning and the end.
- Branch predictor “changes its mind too quickly”

Time	State	Prediction	Outcome	Result?
1	N	N	T	Wrong
2	T	T	T	Correct
3	T	T	T	Correct
4	T	T	N	Wrong
5	N	N	T	Wrong
6	T			
7				
8				
9				
10				
11				
12				

EE156/CS140 Mark Hempstead

14

Branch History Table (BHT)

- **Branch history table (BHT):**
simplest direction predictor

- PC indexes table of bits (0 = N, 1 = T), no tags
- Predicts the branch to go same way it went last time
- Problem: **inner loop branch** below

```
for (i=0; i<100; i++)
  for (j=0; j<3; j++)
    // whatever
```
- It will be wrong twice per inner loop; once each at the beginning and the end.
- Branch predictor “changes its mind too quickly”

Time	State	Prediction	Outcome	Result?
1	N	N	T	Wrong
2	T	T	T	Correct
3	T	T	T	Correct
4	T	T	N	Wrong
5	N	N	T	Wrong
6	T	T	T	Correct
7	T	T	T	Correct
8	T	T	N	Wrong
9	N	N	T	Wrong
10	T	T	T	Correct
11	T	T	T	Correct
12	T	T	N	Wrong

EE156/CS140 Mark Hempstead

15

Two-Bit Saturating Counters (2bc)

• Two-bit saturating counters (2bc)

[Smith 1981]


- Replace each single-bit predictor with
 - $(0,1,2,3) = (N,n,t,T)$
- Adds “hysteresis”
 - Force predictor to mis-predict twice before “changing its mind”
- One mispredict each loop execution (rather than two)
 - + Improves our inner-loop problem.

Time	State	Prediction	Outcome	Result?
1	N	N	T	Wrong
2	n	N	T	Wrong
3				
4				
5				
6				
7				
8				
9				
10				
11				
12				

EE156/CS140 Mark Hempstead

16

Two-bit Saturating Counters



- 2-bit FSMs mean prediction must miss twice before change
- N-bit predictors are possible, but after 2-bits not much benefit

ECEC 621 Mark Hempstead

17

Two-Bit Saturating Counters (2bc)

• Two-bit saturating counters (2bc)

[Smith 1981]

- Replace each single-bit predictor with
 - $(0,1,2,3) = (N,n,t,T)$
- Adds “hysteresis”
 - Force predictor to mis-predict twice before “changing its mind”
- One mispredict each loop execution (rather than two)
 - + Improves our inner-loop problem.

Time	State	Prediction	Outcome	Result?
1	N	N	T	Wrong
2	n	N	T	Wrong
3	t	N	T	Wrong
4				
5				
6				
7				
8				
9				
10				
11				
12				

EE156/CS140 Mark Hempstead

18

Two-Bit Saturating Counters (2bc)

• Two-bit saturating counters (2bc)

[Smith 1981]

- Replace each single-bit predictor with
 - $(0,1,2,3) = (N,n,t,T)$
- Adds “hysteresis”
 - Force predictor to mis-predict twice before “changing its mind”
- One mispredict each loop execution (rather than two)
 - + Improves our inner-loop problem.

Time	State	Prediction	Outcome	Result?
1	N	N	T	Wrong
2	n	N	T	Wrong
3	t	T	T	Correct
4	T	T	T	Correct
5				
6				
7				
8				
9				
10				
11				
12				

EE156/CS140 Mark Hempstead

19

Two-Bit Saturating Counters (2bc)

• Two-bit saturating counters (2bc)

[Smith 1981]

- Replace each single-bit predictor with
 - $(0,1,2,3) = (N,n,t,T)$
- Adds “hysteresis”
 - Force predictor to mis-predict twice before “changing its mind”
- One mispredict each loop execution (rather than two)
 - + Improves our inner-loop problem.

Time	State	Prediction	Outcome	Result?
1	N	N	T	Wrong
2	n	N	T	Wrong
3	t	T	T	Correct
4	T	T	N	Wrong
5	t	T	T	Correct
6				
7				
8				
9				
10				
11				
12				

EE156/CS140 Mark Hempstead

20

Two-Bit Saturating Counters (2bc)

• Two-bit saturating counters (2bc)

[Smith 1981]

- Replace each single-bit predictor with
 - $(0,1,2,3) = (N,n,t,T)$
- Adds “hysteresis”
 - Force predictor to mis-predict twice before “changing its mind”
- One mispredict each loop execution (rather than two)
 - + Improves our inner-loop problem.

Time	State	Prediction	Outcome	Result?
1	N	N	T	Wrong
2	n	N	T	Wrong
3	t	T	T	Correct
4	T	T	N	Wrong
5	t	T	T	Correct
6	T	T	T	Correct
7				
8				
9				
10				
11				
12				

EE156/CS140 Mark Hempstead

21

Two-Bit Saturating Counters (2bc)

• Two-bit saturating counters (2bc)

[Smith 1981]


- Replace each single-bit predictor with
 - (0,1,2,3) = (N,n,t,T)
- Adds “hysteresis”
 - Force predictor to mis-predict twice before “changing its mind”
- One mispredict each loop execution (rather than two)
 - + Improves our inner-loop problem.

Time	State	Prediction	Outcome	Result?
1	N	N	T	Wrong
2	n	N	T	Wrong
3	t	T	T	Correct
4	T	T	N	Wrong
5	t	T	T	Correct
6	T	T	T	Correct
7	T	T	T	Correct
8	T	T	N	Wrong
9	t	T	T	Correct
10	T	T	T	Correct
11	T	T	T	Correct
12	T	T	N	Wrong

EE156/CS140 Mark Hempstead

22

Two-bit Saturating Counters, take 2



- Slightly different version works better for some benchmarks.

ECEC 621 Mark Hempstead

23

Intuition on how well this works

- Simple predictors work well:
 - on branches that repeat the same behavior in streaks.
 - They work better when the streaks are longer
- They don't work well:
 - for random or data-dependent branches.
 - Data dependence can be common in some applications, but it's evil for branch prediction.

EE156/CS140 Mark Hempstead

24

Branch Target Buffer (BTB)


- It's little use predicting taken/not early in ID if you don't know the target address until late in EX.
 - So we record the past branch targets in a hardware structure
- Branch target buffer (BTB):
 - “Last time the branch X was taken, it went to address Y”
 - “So, in the future, if address X is fetched, fetch address Y next”
- Works well because most control insns use **direct targets**
 - Target encoded in insn itself → same “taken” target every time
- What about **indirect targets**?
 - Target held in a register → can be different each time
 - Two indirect call idioms
 - Dynamically linked functions (DLLs): target always the same
 - Dynamically dispatched (virtual) functions: hard but uncommon
 - Also two indirect unconditional jump idioms
 - Switches: hard but uncommon
 - Function returns: hard and common but...

EE156/CS140 Mark Hempstead

25

Implementation


- A small RAM: address = PC, data = target-PC
- Access at Fetch *in parallel* with instruction memory
 - predicted-target = BTB[hash(PC)]
- Updated whenever target != predicted-target
 - BTB[hash(PC)] = correct-target
- Very similar to BHT, but it's a full address width and not just one taken/not bit.
- Aliasing? No problem, this is only a prediction



EE156/CS140 Mark Hempstead

26

Return Address Stack (RAS)




- **Return address stack (RAS)**
 - Call instruction? $RAS[TopOfStack++] = PC+4$
 - Return instruction? Predicted-target = $RAS[--TopOfStack]$
 - Q: how can you tell if an insn is a call/return before decoding it?
 - Accessing RAS on every instruction would waste power
 - Answer:
 - another predictor (or put them in BTB marked as “return”)
 - Or, **pre-decode bits** in insn mem, written when first executed

EE156/CS140 Mark Hempstead

27

Putting It All Together

- BTB & branch direction predictor during fetch



- If branch prediction correct, then no penalty for branches!

EE156/CS140 Mark Hempstead

28

Branch Prediction Performance

- Dynamic branch prediction
 - 20% of instruction branches
 - Simple predictor: branches predicted with 75% accuracy
 - $CPI = 1 + (20\% * 25\% * 2) = 1.1$
 - More advanced predictor: 95% accuracy
 - $CPI = 1 + (20\% * 5\% * 2) = 1.02$
- Branch mis-predictions still a big problem though
 - Pipelines are long: typical mis-prediction penalty is 10+ cycles
 - For cores that do more per cycle, predictions most costly (later)
- Many other methods for building a better branch predictor including correlating predictors, global history tables (Chapter 3)


EE156/CS140 Mark Hempstead

29

BHT Accuracy

- Mispredict because either:
 - Wrong guess for that branch
 - Got branch history of wrong branch due to aliasing (4K-entry table shown here).

Another way:
Prediction
accuracy between
99% and 82%



ECEC 621 Mark Hempstead

30

ADVANCED BRANCH PREDICTION (3.3)

ECEC 621
Mark Hempstead

31

Many more effective and complex branch prediction techniques exist

- Correlating branch predictors (sec 3.3)
 - Gshare correlating predictor (paper and lab)
- Tournament predictors (sec 3.3)
- Tagged Hybrid predictors (sec 3.3)
- Core i7 branch predictors (sec 3.3)
- Predictors base on machine learning
 - Perceptron branch prediction (see paper)

EE156/CS140 Mark Hempstead

32

Correlating Predictors

- 2-bit scheme only looks at branch's *own* history to predict its behavior
- What if we use other branches to predict it as well?

```
if (aa==2) aa=0; // Branch #1
if (bb==2) bb=0; // Branch #2
if (aa!=bb) { . . } // Branch #3
```

- Does branch #3 depends on outcome of #1 and #2?
 - Yes. If #1 & #2 are both taken, then #3 will not be taken.
 - Sometimes the past can predict the future with certainty!
 - This is called “global history.”

ECEC 621 Mark Hempstead

33

Another example of global history

- Consider our inner loop again
 - The branch pattern is (T,T,T,N) repeating
- Consider the following global rule:
 - If the last branches were "T,T,T" then predict not taken
 - Else predict taken.
 - Works perfectly – the trick is having the predictor figure out that rule.
- Let's make a predictor that looks at both the history at this branch, as well as recent outcomes of all branches (including this one).

EE156/CS140 Mark Hempstead

34


Branch History Register

- Simple shift register
 - Shift in branch outcomes as they occur
 - 1 => branch was taken
 - 0 => branch was not-taken
 - k-bit BHR => 2^k patterns
 - Now we know the outcome of the last k branches.
 - Use these patterns to address into the Pattern History Table. Effectively the PHT is a separate BHT for each global pattern.

ECEC 621 Mark Hempstead

35

Generic Adaptive Branch Prediction (Correlating Predictor)



- Two-level BP requires two main components
 - Branch history register (BHR): recent outcomes of branches (last k branches encountered)
 - Pattern History Table (PHT): branch behavior for recent occurrences of the specific pattern of these k branches
 - In effect, we concatenate BHR with Branch PC bits
- We've drawn a (2,2) predictor

ECEC 621 Mark Hempstead

36

Pattern History Table

- Has 2^k entries
- Usually uses a 2-bit counter for the prediction
- BHR is used to address the PHT. I.e., a $(k,2)$ PHT is an array of 2^k BHTs, each of which has a 2-bit counter.

ECEC 621 Mark Hempstead

37

(2,1) correlated Predictor

• Correlated (two-level)

predictor [Patt 1991]

- Exploits observation that branch outcomes are correlated
- Maintains separate prediction per (PBHR, PC) pairs
 - **Branch history register (BHR):** recent branch outcomes
- Simple working example: assume program has one branch
 - BHT: one 1-bit DIRP entry
 - BHT+2BHR: $2^2 = 4$ 1-bit DIRP entries

Time	BHR	PHT				Prediction	Outcome	Result?
		NN	NT	NT	TT			
1	NN	N	N	N	N		T	
2							T	
3							T	
4							N	
5							T	
6							T	
7							T	
8							N	
9							T	
10							T	
11							T	
12							N	

38

(2,1) correlated Predictor

• Correlated (two-level)

predictor [Patt 1991]

- Exploits observation that branch outcomes are correlated
- Maintains separate prediction per (PC, BHR) pairs
 - **Branch history register (BHR):** recent branch outcomes
- Simple working example: assume program has one branch
 - BHT: one 1-bit DIRP entry
 - BHT+2BHR: $2^2 = 4$ 1-bit DIRP entries

Time	BHR	PHT				Prediction	Outcome	Result?
		NN	NT	NT	TT			
1	NN	N	N	N	N		T	
2	NT	N	N	N	N		T	Wrong
3								
4								
5								
6								
7								
8								
9								
10								
11								
12								

39

(2,1) correlated Predictor

• Correlated (two-level) predictor [Patt 1991]

- Exploits observation that branch outcomes are correlated
- Maintains separate prediction per (PC, BHR) pairs

- **Branch history register (BHR):** recent branch outcomes

- Simple working example: assume program has one branch

- BHT: one 1-bit DIRP entry
- BHT+2BHR: $2^2 = 4$ 1-bit DIRP entries

Time	BHR	PHT				Prediction	Outcome	Result?
		NN	NT	NT	TT			
1	NN	N	N	N	N	N	T	Wrong
2	NT	T	N	N	N	N	T	Wrong
3	TT	T	T	N	N	N	T	Wrong
4	TT	T	T	N	T	N	T	Wrong
5								
6								
7								
8								
9								
10								
11								
12								

40

(2,1) correlated Predictor

• Correlated (two-level) predictor [Patt 1991]

- Exploits observation that branch outcomes are correlated
- Maintains separate prediction per (PC, BHR) pairs

- **Branch history register (BHR):** recent branch outcomes

- Simple working example: assume program has one branch

- BHT: one 1-bit DIRP entry
- BHT+2BHR: $2^2 = 4$ 1-bit DIRP entries

Time	BHR	PHT				Prediction	Outcome	Result?
		NN	NT	NT	TT			
1	NN	N	N	N	N	N	T	Wrong
2	NT	T	N	N	N	N	T	Wrong
3	TT	T	T	N	N	N	T	Wrong
4	TT	T	T	N	T	N	T	Wrong
5								
6								
7								
8								
9								
10								
11								
12								

41

(2,1) correlated Predictor

• Correlated (two-level) predictor [Patt 1991]

- Exploits observation that branch outcomes are correlated
- Maintains separate prediction per (PC, BHR) pairs

- **Branch history register (BHR):** recent branch outcomes

- Simple working example: assume program has one branch

- BHT: one 1-bit DIRP entry
- BHT+2BHR: $2^2 = 4$ 1-bit DIRP entries

Time	BHR	PHT				Prediction	Outcome	Result?
		NN	NT	NT	TT			
1	NN	N	N	N	N	N	T	Wrong
2	NT	T	N	N	N	N	T	Wrong
3	TT	T	T	N	N	N	T	Wrong
4	TT	T	T	N	T	N	T	Wrong
5	TN	T	T	N	N	N	T	Wrong
6								
7								
8								
9								
10								
11								
12								

42

(2,1) correlated Predictor

• Correlated (two-level) predictor [Patt 1991]

- Exploits observation that branch outcomes are correlated
- Maintains separate prediction per (PC, BHR) pairs

- **Branch history register (BHR):** recent branch outcomes

- Simple working example: assume program has one branch

- BHT: one 1-bit DIRP entry
- BHT+2BHR: $2^2 = 4$ 1-bit DIRP entries

Time	BHR	PHT				Prediction	Outcome	Result?
		NN	NT	NT	TT			
1	NN	N	N	N	N	N	T	Wrong
2	NT	T	N	N	N	N	T	Wrong
3	TT	T	T	N	N	N	T	Wrong
4	TT	T	T	N	T	T	N	Wrong
5	TN	T	T	N	N	N	T	Wrong
6	NT	T	T	T	N			
7								
8								
9								
10								
11								
12								

43

(2,1) correlated Predictor

• Correlated (two-level) predictor [Patt 1991]

- Exploits observation that branch outcomes are correlated
- Maintains separate prediction per (PC, BHR) pairs

- **Branch history register (BHR):** recent branch outcomes

- Simple working example: assume program has one branch

- BHT: one 1-bit DIRP entry
- BHT+2BHR: $2^2 = 4$ 1-bit DIRP entries

Time	BHR	PHT				Prediction	Outcome	Result?
		NN	NT	NT	TT			
1	NN	N	N	N	N	N	T	Wrong
2	NT	T	N	N	N	N	T	Wrong
3	TT	T	T	N	N	N	T	Wrong
4	TT	T	T	N	T	T	N	Wrong
5	TN	T	T	N	N	N	T	Wrong
6	NT	T	T	T	N	T	T	Correct
7	TT	T	T	T	N			
8								
9								
10								
11								
12								

44

(2,1) correlated Predictor

• Correlated (two-level) predictor [Patt 1991]

- Exploits observation that branch outcomes are correlated
- Maintains separate prediction per (PC, BHR) pairs

- **Branch history register (BHR):** recent branch outcomes

- Simple working example: assume program has one branch

- BHT: one 1-bit DIRP entry
- BHT+2BHR: $2^2 = 4$ 1-bit DIRP entries

Time	BHR	PHT				Prediction	Outcome	Result?
		NN	NT	NT	TT			
1	NN	N	N	N	N	N	T	Wrong
2	NT	T	N	N	N	N	T	Wrong
3	TT	T	T	N	N	N	T	Wrong
4	TT	T	T	N	T	T	N	Wrong
5	TN	T	T	N	N	N	T	Wrong
6	NT	T	T	T	N	T	T	Correct
7	TT	T	T	T	N	N	T	Wrong
8	TT	T	T	T	T			
9								
10								
11								
12								

45

(2,1) correlated Predictor

• Correlated (two-level) predictor [Patt 1991]

- Exploits observation that branch outcomes are correlated
- Maintains separate prediction per (PC, BHR) pairs

- **Branch history register (BHR):** recent branch outcomes

- Simple working example: assume program has one branch

- BHT: one 1-bit DIRP entry
- BHT+2BHR: $2^2 = 4$ 1-bit DIRP entries

Time	BHR	PHT				Prediction	Outcome	Result?
		NN	NT	TT	TT			
1	NN	N	N	N	N	N	T	Wrong
2	NT	T	N	N	N	N	T	Wrong
3	TT	T	T	N	N	N	T	Wrong
4	TT	T	T	N	T	T	N	Wrong
5	TN	T	T	N	N	N	T	Wrong
6	NT	T	T	T	N	T	T	Correct
7	TT	T	T	T	N	N	T	Wrong
8	TT	T	T	T	T	T	N	Wrong
9	TN	T	T	T	N	T	T	Correct
10								
11								
12								

46

(2,1) correlated Predictor

• Correlated (two-level) predictor [Patt 1991]

- Exploits observation that branch outcomes are correlated
- Maintains separate prediction per (PC, BHR) pairs

- **Branch history register (BHR):** recent branch outcomes

- Simple working example: assume program has one branch

- BHT: one 1-bit DIRP entry
- BHT+2BHR: $2^2 = 4$ 1-bit DIRP entries

Time	BHR	PHT				Prediction	Outcome	Result?
		NN	NT	TT	TT			
1	NN	N	N	N	N	N	T	Wrong
2	NT	T	N	N	N	N	T	Wrong
3	TT	T	T	N	N	N	T	Wrong
4	TT	T	T	N	T	T	N	Wrong
5	TN	T	T	N	N	N	T	Wrong
6	NT	T	T	T	N	T	T	Correct
7	TT	T	T	T	N	N	T	Wrong
8	TT	T	T	T	T	T	N	Wrong
9	TN	T	T	T	N	T	T	Correct
10	NT	T	T	T	N	T	T	Correct
11								
12								

47

(2,1) correlated Predictor

• Correlated (two-level) predictor [Patt 1991]

- Exploits observation that branch outcomes are correlated
- Maintains separate prediction per (PC, BHR) pairs

- **Branch history register (BHR):** recent branch outcomes

- Simple working example: assume program has one branch

- BHT: one 1-bit DIRP entry
- BHT+2BHR: $2^2 = 4$ 1-bit DIRP entries

Time	BHR	PHT				Prediction	Outcome	Result?
		NN	NT	TT	TT			
1	NN	N	N	N	N	N	T	Wrong
2	NT	T	N	N	N	N	T	Wrong
3	TT	T	T	N	N	N	T	Wrong
4	TT	T	T	N	T	T	N	Wrong
5	TN	T	T	N	N	N	T	Wrong
6	NT	T	T	T	N	T	T	Correct
7	TT	T	T	T	N	N	T	Wrong
8	TT	T	T	T	T	T	N	Wrong
9	TN	T	T	T	N	T	T	Correct
10	NT	T	T	T	N	T	T	Correct
11	TT	T	T	T	N	T	T	Correct
12								

48

(2,1) correlated Predictor

• Correlated (two-level) predictor [Patt 1991]

- Exploits observation that branch outcomes are correlated
- Maintains separate prediction per (PC, BHR) pairs
 - Branch history register (BHR):** recent branch outcomes
- Simple working example: assume program has one branch
 - BHT: one 1-bit DIRP entry
 - BHT+2BHR: $2^2 = 4$ 1-bit DIRP entries

Time	BHR	PHT				Prediction	Outcome	Result?
		NN	NT	TN	TT			
1	NN	N	N	N	N	N	T	Wrong
2	NT	T	N	N	N	N	T	Wrong
3	TT	T	T	N	N	N	T	Wrong
4	TT	T	T	N	T	T	N	Wrong
5	TN	T	T	N	N	N	T	Wrong
6	NT	T	T	T	N	T	T	Correct
7	TT	T	T	T	N	N	T	Wrong
8	TT	T	T	T	T	T	N	Wrong
9	TN	T	T	T	N	T	T	Correct
10	NT	T	T	T	N	T	T	Correct
11	TT	T	T	T	N	N	T	Wrong
12	TT	T	T	T	T	T	N	Wrong

49

(2,1) correlated Predictor

• What went wrong?

- Look at the grey squares; they log our mistakes.
- For NN, NT, TN we only made one mistake. But TT kept making mistakes. Why?
 - We have both T,T,T and also T,T,N.
 - Every time we went from one to the other we got the wrong answer and had to re-train our predictor.
- How do we fix this simply?
 - Go to a (3,1) predictor.

Time	BHR	PHT				Prediction	Outcome	Result?
		NN	NT	TN	TT			
1	NN	N	N	N	N	N	T	Wrong
2	NT	T	N	N	N	N	T	Wrong
3	TT	T	T	N	N	N	T	Wrong
4	TT	T	T	N	T	T	N	Wrong
5	TN	T	T	N	N	N	T	Wrong
6	NT	T	T	T	N	T	T	Correct
7	TT	T	T	T	N	N	T	Wrong
8	TT	T	T	T	T	T	N	Wrong
9	TN	T	T	T	N	T	T	Correct
10	NT	T	T	T	N	T	T	Correct
11	TT	T	T	T	N	N	T	Wrong
12	TT	T	T	T	T	T	N	Wrong

50

(3,1) correlated predictor

- 3-bit BHR
- 2^3 DIRP entries per pattern
- One-bit BHT (not a 2-bit saturating counter)

Time	BHR	PHT								Prediction	Outcome	Result?
		NNN	NNT	NTN	NTT	TNN	TNT	TTN	TTT			
1												
2												
3												
4												
5												
6												
7												
8												
9												
10												
11												
12												

51

(3,1) correlated predictor

- 3-bit BHR
- 2^3 DIRP entries per pattern
- One-bit BHT (not a 2-bit saturating counter)

Time	BHR	PHT								Prediction	Outcome	Result?
		NNN	NNT	NTN	NTT	TNN	TNT	TTN	TTT			
1	NNN	N	N	N	N	N	N	N	N	N	T	Wrong
2	NNT	T	N	N	N	N	N	N	N			
3	NTT											
4	TTT											
5												
6												
7												
8												
9												
10												
11												
12												

52

(3,1) correlated predictor

- 3-bit BHR
- 2^3 DIRP entries per pattern
- One-bit BHT (not a 2-bit saturating counter)

Time	BHR	PHT								Prediction	Outcome	Result?
		NNN	NNT	NTN	NTT	TNN	TNT	TTN	TTT			
1	NNN	N	N	N	N	N	N	N	N	N	T	Wrong
2	NNT	T	N	N	N	N	N	N	N	N	T	Wrong
3	NTT	T	T	N	N	N	N	N	N			
4	TTT											
5												
6												
7												
8												
9												
10												
11												
12												

53

(3,1) correlated predictor

- 3-bit BHR
- 2^3 DIRP entries per pattern
- One-bit BHT (not a 2-bit saturating counter)

Time	BHR	PHT								Prediction	Outcome	Result?
		NNN	NNT	NTN	NTT	TNN	TNT	TTN	TTT			
1	NNN	N	N	N	N	N	N	N	N	N	T	Wrong
2	NNT	T	N	N	N	N	N	N	N	N	T	Wrong
3	NTT	T	T	N	N	N	N	N	N	N	T	Wrong
4	TTT	T	T	N	T	N	N	N	N			
5												
6												
7												
8												
9												
10												
11												
12												

54

(3,1) correlated predictor

- 3-bit BHR
- 2³ DIRP entries per pattern
- One-bit BHT (not a 2-bit saturating counter)

Time	BHR	PHT								Prediction	Outcome	Result?
		NNN	NNT	NTN	NTT	TNN	TNT	TTN	TTT			
1	NNN	N	N	N	N	N	N	N	N	N	T	Wrong
2	NNT	T	N	N	N	N	N	N	N	N	T	Wrong
3	NTT	T	T	N	N	N	N	N	N	N	T	Wrong
4	TTT	T	T	N	T	N	N	N	N	N	N	Correct
5	TTN	T	T	N	T	N	N	N	N	N	T	Wrong
6												
7												
8												
9												
10												
11												
12												

55

(3,1) correlated predictor

- 3-bit BHR
- 2³ DIRP entries per pattern
- One-bit BHT (not a 2-bit saturating counter)

Time	BHR	PHT								Prediction	Outcome	Result?
		NNN	NNT	NTN	NTT	TNN	TNT	TTN	TTT			
1	NNN	N	N	N	N	N	N	N	N	N	T	Wrong
2	NNT	T	N	N	N	N	N	N	N	N	T	Wrong
3	NTT	T	T	N	N	N	N	N	N	N	T	Wrong
4	TTT	T	T	N	T	N	N	N	N	N	N	Correct
5	TTN	T	T	N	T	N	N	N	N	N	T	Wrong
6	TNT	T	T	N	T	N	N	N	N	N	T	Wrong
7												
8												
9												
10												
11												
12												

56

(3,1) correlated predictor

- 3-bit BHR
- 2³ DIRP entries per pattern
- One-bit BHT (not a 2-bit saturating counter)

Time	BHR	PHT								Prediction	Outcome	Result?
		NNN	NNT	NTN	NTT	TNN	TNT	TTN	TTT			
1	NNN	N	N	N	N	N	N	N	N	N	T	Wrong
2	NNT	T	N	N	N	N	N	N	N	N	T	Wrong
3	NTT	T	T	N	N	N	N	N	N	N	T	Wrong
4	TTT	T	T	N	T	N	N	N	N	N	N	Correct
5	TTN	T	T	N	T	N	N	N	N	N	T	Wrong
6	TNT	T	T	N	T	N	N	N	N	N	T	Wrong
7	NTT	T	T	N	T	N	N	N	N	N	T	Wrong
8												
9												
10												
11												
12												

57

(3,1) correlated predictor

- 3-bit BHR
- 2^3 DIRP entries per pattern
- One-bit BHT (not a 2-bit saturating counter)

Time	BHR	PHT								Prediction	Outcome	Result?
		NNN	NNT	NTN	NTT	TNN	TNT	TTN	TTT			
1	NNN	N	N	N	N	N	N	N	N	N	T	Wrong
2	NNT	T	N	N	N	N	N	N	N	N	T	Wrong
3	NTT	T	T	N	N	N	N	N	N	N	T	Wrong
4	TTT	T	T	N	T	N	N	N	N	N	N	Correct
5	TTN	T	T	N	T	N	N	N	N	N	T	Wrong
6	TNT	T	T	N	T	N	N	N	N	N	T	Wrong
7	NTT	T	T	N	T	N	T	T	N	T	T	Correct
8	TTT	T	T	N	T	N	T	T	N	T	T	Correct
9												
10												
11												
12												

58

(3,1) correlated predictor

- With 3 bits of history, the history is now enough to fully predict the branch result.
- Any column has only one grey square; once we train the predictor for a given history it never makes another mistake.

Time	BHR	PHT								Prediction	Outcome	Result?
		NNN	NNT	NTN	NTT	TNN	TNT	TTN	TTT			
1	NNN	N	N	N	N	N	N	N	N	N	T	Wrong
2	NNT	T	N	N	N	N	N	N	N	N	T	Wrong
3	NTT	T	T	N	N	N	N	N	N	N	T	Wrong
4	TTT	T	T	N	T	N	N	N	N	N	N	Correct
5	TTN	T	T	N	T	N	N	N	N	N	T	Wrong
6	TNT	T	T	N	T	N	N	N	N	N	T	Wrong
7	NTT	T	T	N	T	N	T	T	N	T	T	Correct
8	TTT	T	T	N	T	N	T	T	N	N	N	Correct
9	TTN	T	T	N	T	N	T	T	N	T	T	Correct
10	TNT	T	T	N	T	N	T	T	N	T	T	Correct
11	NTT	T	T	N	T	N	T	T	N	T	T	Correct
12	TTT	T	T	N	T	N	T	T	N	N	N	Correct


59

Hardware Costs of 2-level predictions

- (m,n) predictor → m-bits of global history, n-bit predictor
- $2^m \cdot n$ * Number of prediction entries
- Say you have m-bits of history (m=2)
- n-bits of predictor per entries (n=2)

(2,2) predictor with 1K prediction entries
 $2^2 \cdot 2 \cdot 1024 = 8K$ -bits

Accuracy of Different Schemes




ECEC 621 Mark Hempstead

61

Hybrid Branch Predictors

- Tournament predictors: Adaptively combine local and global predictors
- Different schemes work better for different branches



ECEC 621 Mark Hempstead

62

Tagged Hybrid Predictors

- Need methods that combine multiple predictors and then evaluate if the prediction is likely to be associated with the branch.
- This association can depend on branch histories of varying length.
- According to the textbook it is believed that Intel and others are employing this predictor style in their microprocessors. Though it has not been disclosed publically.

EE156/CS140 Mark Hempstead

63


Tagged Hybrid Predictors

- Need to have predictor for each branch and history
 - Problem: this implies huge tables
 - Solution:
 - Use hash tables, whose hash value is based on branch address and branch history
 - Longer histories may lead to increased chance of hash collision, so use multiple tables with increasingly shorter histories

64

Tagged Hybrid Predictors


- Example 5 component predictor – 5 predictor tables
- Accessed with a hash of the PC and history of last i branches (like in gshare)
- Each table has different history lengths
- Use of a short tag to determine a match
- Use field indicates if the prediction was recently used and helps with match
- Typically 32-64KiB in size



- Problems: how to initialize, need to regularly reset history
- Could use hints from software

65


Tagged Hybrid Predictors



66

Other Branch Predictors

- There are many branch predictors developed over the years.
- We will read:
 - Perceptron Branch Predictor
- We will implement
 - gshare



EE156/CS140 Mark Hempstead

67

Branch prediction costs

- The BTB and BHT are big arrays that run every single cycle.
 - They are on critical paths, and thus must run fast.
 - They can be big power hogs.
 - They will never work well if there are data-dependent branches.
 - For long pipes at high frequencies, a good branch prediction is very necessary.
 - For shorter pipes (where the mis-predict penalty is less severe), you can get away with a less elaborate predictor (e.g., smaller arrays).

EE156/CS140 Mark Hempstead

68

More costs

- “When you come to a fork in the road, take it”
 - Yogi Berra
- The central idea of branch prediction is that doing nothing is bad.
 - Don’t stall. Instead, *speculatively execute* a (hopefully pretty good) prediction
 - But executing instructions takes energy, and if you have to flush them then the energy was wasted.
- Unless you’re predictors are *always* right, then branch prediction + speculative execution *always* wastes energy in flushed instructions.
- Nonetheless, the tradeoff is *usually* a good one.

EE156/CS140 Mark Hempstead

69